

Milestone 4 – Client Acceptance Testing Guide

Stakeholder URL: <http://52.70.0.80>

Milestone: Output Generation, Testing & Deployment

Build tag: `m4-20260627-191552-ed97574` (app + worker)

Structured output schema: `trueblue.chat.output.v1`

Prerequisite: Milestone 3 (vector database & grounded AI chat) is available

Important – read before uploading anything: The current stakeholder environment is served over **HTTP, not HTTPS**, and the application does **not yet redact PII** (SSN, bank/account numbers, email, phone) from answers. For M4 acceptance testing, use **only non-sensitive or redacted sample PDFs**. Do **not** upload real client tax documents until TLS/HTTPS and an output-layer redaction pass are complete or the risk is explicitly waived in writing. Open <http://52.70.0.80/login> directly with `http://`, not `https://`.

Contents

- [What This Milestone Proves](#)
 - [What M4 Delivers \(Mapped to Scope\)](#)
 - [Access](#)
 - [Test Environment](#)
 - [Before You Start](#)
 - [UAT Test Matrix](#)
 - [Detailed Test Walkthroughs](#)
 - [Scope and Limitations](#)
 - [Acceptance Criteria Checklist](#)
 - [Sign-off](#)
-

What This Milestone Proves

Milestone 4 validates the complete document Q&A path from upload through a **structured, source-cited response**, delivered through a chat interface on a deployed staging environment.

Plain-language note – "structured envelope": alongside the readable answer, each response also carries a **structured envelope** – a single machine-readable record that bundles the answer text together with its source citations, page references, and support details. M4 proves this envelope is produced correctly; it is what you inspect in the **Raw / Structured JSON** view.

Scope acceptance criteria being tested (from the Phase 1 Statement of Work):

1. **End-to-end flow is functional:** upload a PDF, query it, and receive a **structured JSON response with source citations**.
2. **Output conforms to the agreed JSON schema** (`trueblue.chat.output.v1`).
3. **Chat interface allows query submission and displays structured responses with source citations**.
4. **System is deployed and accessible at the agreed endpoint/URL**.
5. **Deployment documentation is sufficient for an engineer to redeploy independently**.

What M4 deliberately does **not** cover (so testing stays in scope):

- OCR for scanned/image-only PDFs – **Milestone 5**.
- The full tax-domain **8-section advisory report**, tax-strategy categorization, advice disclaimers, usage dashboards, and admin rate-limit controls – **Milestone 6**. M4 proves the *grounded, source-cited structured envelope*, not advisory generation.

What M4 Delivers (Mapped to Scope)

SOW deliverable	What was built	Where to verify
Structured JSON output format (response text, source references, page/section metadata, confidence/relevance)	A server-built <code>trueblue.chat.output.v1</code> envelope: <code>responseText</code> , <code>sources[]</code> , <code>coverage</code> , <code>support</code> (confidence label + basis + retrieval mode), <code>warnings</code> , <code>metadata</code> .	TC-02
Configurable output template engine with support for client-defined schemas	Registered templates <code>rag_qa.default.v1</code> (default) and <code>rag_qa.compact.v1</code> . Unknown templates are rejected. Client-defined schemas are supported as approved/registered templates, not arbitrary runtime JSON schemas (see Limitation 2).	TC-05 (confirmed by developer)
Chat/query interface with citations	A full-bleed chat workspace: query input, streamed answer with inline <code>[S1]</code> markers, source selection, citation detail (filename + page + snippet), an Insights view and a Raw / Structured JSON view, and a per-answer action bar (Copy / Regenerate).	TC-06, TC-07, TC-08

SOW deliverable	What was built	Where to verify
End-to-end pipeline (upload → process → embed → query → response)	The full pipeline runs on staging and is exercised by the browser walkthroughs below.	TC-01
Deployment to agreed staging environment	Deployed to http://52.70.0.80 (AWS account 536573256060), build tag m4-20260627-191552-ed97574 .	TC-10
Deployment documentation	Deployment notes, the output-contract reference, and this acceptance guide, delivered to the developer for independent redeploy.	TC-11 (developer-attested)

Access

URL: <http://52.70.0.80> (HTTP only – see the caveat at the top of this document).

Test accounts (staging seed)

Email	Password	Role	Firm
admin@acmetax.com	FirmAdmin1!	Firm Admin	Acme Tax Services
user@acmetax.com	FirmUser1!	Firm User	Acme Tax Services
admin@besttax.com	FirmAdmin1!	Firm Admin	Best Tax Advisors
admin@trueblue.dev	Admin123!	Platform Admin	All firms

Recommended account for most browser checks:

`admin@acmetax.com / FirmAdmin1!`

Use `admin@besttax.com` only for the tenant-isolation check (TC-09). The Platform Admin account is **not** the primary document Q&A account, because M4 chat and upload are firm-scoped.

Registration firm codes (if you create your own user): `acme-tax` , `best-tax` .

These are **throwaway** staging passwords on an HTTP environment. Do not reuse them anywhere real.

Sample PDFs

Use the same non-sensitive sample tax returns used for M2/M3. The pre-verified sample document for this guide is:

- [2025 Tax Return Documents \(Whittaker Jordan\).pdf](#) (25 pages)— used for the worked examples in TC-01 and TC-08 below.

Other good candidates for the multi-source check:

- [2025 Tax Return Documents \(Jimenez Julio\).pdf](#)
- [2025 Tax Return Documents \(SMITH TALIA S and Antonio Smith\).pdf](#)

At least one text-based PDF must be uploaded or already present with status `COMPLETED` .

Test Environment

Property	Value
Endpoint	<code>http://52.70.0.80</code> (HTTP only)
Build tag	<code>m4-20260627-191552-ed97574</code> (app + worker)
AWS account	<code>536573256060</code> , region <code>us-east-1</code>
Schema version	<code>trueblue.chat.output.v1</code>


If the `Sources` area shows no completed documents, upload one sample PDF and wait for `COMPLETED` before continuing.

Before You Start

1. Open <http://52.70.0.80/login> (use `http://`).
2. Log in as `admin@acmetax.com` .
3. Open <http://52.70.0.80/dashboard/chat>.
4. Confirm the chat workspace loads: a query/composer area, a way to select **Sources**, a conversation area, and (after an answer) an **Insights** view and a **Raw / Structured JSON** view.

Finding your way around

The tests below reference a few on-screen controls. Here is where each one lives:



- **Sources** panel — lists your firm's completed documents and lets you tick the ones a question should use. Its **upload / "Choose files"** control is how you add a new PDF.
- **History (Chats)** list — your saved conversations. Click a row to reopen a past chat; use **New thread** to start a fresh one.
- **Suggested-prompt chips** — the one-click example questions shown in an empty chat; you can click one or type your own question in the composer.
- **Answer + citations** — the answer streams into the conversation area. Cited passages show an inline marker such as `[S1]` ; expand a citation to see the source filename, page, and snippet.
- **Insights vs Raw / Structured JSON** views — once an answer appears, a small view switch beside the answer lets you toggle between **Insights** (a plain-language summary: status, support, coverage) and **Raw / Structured JSON** (the full structured envelope as JSON). Several tests ask you to switch to the **Raw / Structured JSON** view.

Sequential uploads: uploads are processed **one at a time per app process**. If two testers upload simultaneously, one may see `429 Too Many Requests` . Upload **one document at a time** and wait for `COMPLETED` before the next.

Upload size: keep test PDFs under **20 MB**. The UI accepts files up to 25 MB, but the API rejects files over 20 MB (see Limitation 4). All provided sample PDFs are well under 20 MB.

Reporting issues during UAT: record results in the [Acceptance Criteria Checklist](#) below. If a test fails or a result is unclear, **report it to the developer** with the question you asked, the response you got, and a screenshot.

UAT Test Matrix

Status legend: **Client-to-run** = you execute this during UAT and record PASS/FAIL. **Developer-run / Developer-attested** = confirmed by the developer and reported; you acknowledge it in Sign-off (no client action required).

TC	Scope criterion / deliverable	Steps (summary)	Expected result	Run by
TC-01	C1 — E2E upload → process → query → structured JSON w/ citations	Upload (or select) the Whittaker text PDF; wait for <code>COMPLETED</code> ; select it; ask the filing-status / wages question.	Grounded answer (filing status = Single; wages / total income = \$27,645) with an inline <code>[S1]</code> citation whose cited source filename matches the document you selected ; the Raw / Structured JSON view shows <code>output.schemaVersion = trueblue.chat.output.v1</code> , <code>status = answered</code> , and <code>support.retrievalMode = vector_retrieval</code> .	Client-to-run

TC	Scope criterion / deliverable	Steps (summary)	Expected result	Run by
TC-02	C2 – Output conforms to <code>trueblue.chat.output.v1</code>	Open the Raw / Structured JSON view for an answered response.	All required fields present: <code>schemaVersion</code> , <code>templateId</code> , <code>templateVersion</code> , <code>status</code> , <code>responseText</code> , <code>sources[]</code> , <code>coverage</code> , <code>support</code> , <code>warnings</code> , <code>metadata.responseText</code> matches the visible answer; <code>sources[]</code> are final cited sources only with page metadata.	Client-to-run
TC-03	C1/C2 – Insufficient-evidence behavior	With a source selected, ask a question the document cannot answer (e.g. "What spacecraft purchase is in this return?").	Assistant does not invent an answer; <code>status = insufficient_evidence</code> ; <code>sources = []</code> ; warning <code>INSUFFICIENT_EVIDENCE</code> ; no citations shown.	Client-to-run
TC-04	C3 – Non-document message hygiene	Ask <code>hi</code> (a greeting) with a source selected.	Brief guidance, no citations; <code>status = non_document</code> ; <code>sources = []</code> .	Client-to-run
TC-05	Deliverable – configurable output templates / "client-defined schemas"	Developer-run: the developer requests the compact registered template (<code>rag_qa.compact.v1</code>), then an unregistered template id, and reports the result.	Compact request returns <code>templateId = rag_qa.compact.v1</code> with valid cited sources. Unknown template returns HTTP 400 (rejected). Confirms client-defined schemas are supported as approved registered templates , not arbitrary runtime schemas.	Developer-run (confirmed by developer)
TC-06	C3 – Chat UI shows structured responses + citations	In <code>/dashboard/chat</code> , ask a supported question; expand the citation and the structured views.	Answer streams into the conversation; inline <code>[S1]</code> marker; citation expands to filename + page + snippet; Insights view shows status/support/coverage; Raw / Structured JSON view shows the full <code>output</code> object; cited sources are marked "Used".	Client-to-run
TC-07	C3 – Thread persistence & replay	Ask a supported question; confirm a History (Chats) entry; reopen it.	Previous messages reload; the saved answer reopens with its citations and structured details intact ; the locked source scope is preserved.	Client-to-run

TC	Scope criterion / deliverable	Steps (summary)	Expected result	Run by
TC-08	C3 – Source selection + multi-source coverage honesty	Start a new thread; select 2–3 completed PDFs (including the Whittaker return); ask "For each selected return, what taxpayer name is shown?"	Each selected return is addressed (Whittaker return → Jordan Whittaker); a return that lacks evidence for the field is explicitly called out (not silently dropped); citations correspond to the documents actually used; <i>coverage</i> reflects selected vs cited documents.	Client-to-run
TC-09	Security spot-check – tenant isolation	Confirm Acme docs/chats as <i>admin@acmetax.com</i> ; sign out; log in as <i>admin@besttax.com</i> ; confirm Acme data is not visible.	Best Tax cannot see Acme documents or chats. (See Limitation 10 for the deferred prompt-premise wording item.)	Client-to-run
TC-10	C4 – Deployed and accessible	Open <i>http://52.70.0.80/login</i> ; log in; open <i>/dashboard/chat</i> .	Login page loads; login succeeds; dashboard and chat load without server errors at the agreed endpoint.	Client-to-run
TC-11	C5 – Deployment documentation redeployable	No client action. The developer attests the deployment documentation delivered with this handoff is sufficient to redeploy and re-run M4 validation independently.	Developer confirms the deployment docs are sufficient; you acknowledge this in Sign-off.	Developer-attested

Detailed Test Walkthroughs

The matrix above is the spine of sign-off. The walkthroughs below give step-by-step detail for the browser checks.

TC-01 – End-to-end: upload → process → query → structured response

1. Open `http://52.70.0.80/dashboard/chat`.
2. In the **Sources** panel, upload a non-sensitive text-based sample PDF (the **Whittaker** return is the pre-verified sample) using the **upload / "Choose files"** control – one at a time (see the sequential-upload note).
3. Wait for the document to reach `COMPLETED`.
4. Tick the document in **Sources** and ask:

What filing status and wages appear in this return?

Expected (PASS):

- The answer is grounded in the selected document. For the Whittaker return, **expect filing status = Single** and **wages / total income = \$27,645**, and the answer cites a page of the Whittaker return with an inline marker such as `[S1]`.
- The **cited source filename matches the document you selected**.
- Switch to the **Raw / Structured JSON** view and confirm:
 - `schemaVersion = trueblue.chat.output.v1`
 - `status = answered`
 - `support.retrievalMode = vector_retrieval` – this means the AI searched the document index for the most relevant passages. **If it instead shows `local_retrieval_fallback`, the document search index was not used – that is a FAIL; report it to the developer.**



Grounded answer with citations

Wording may vary slightly between runs, but the cited values and the source must match the Whittaker return.

TC-02 – Output conforms to `trueblue.chat.output.v1`

How to open the Raw / Structured JSON view: after an answer appears, use the small view switch beside the answer (the **Insights / Raw** toggle) and select **Raw / Structured JSON**. The full structured envelope appears as JSON.

Confirm every field below is present, and that `support.retrievalMode` reads `vector_retrieval`:

Field	Expected
<code>schemaVersion</code>	Exactly <code>trueblue.chat.output.v1</code>

Field	Expected
templateId / templateVersion	rag_qa.default.v1 / 1
status	answered for a supported question (other values: insufficient_evidence , non_document)
responseText	Matches the visible answer
sources[]	Final cited sources only (marker/sourceId, documentId, page metadata, snippet)
coverage	Selected / retrieved / final / no-evidence document coverage
support	confidenceLabel , confidenceBasis , retrievalMode , source counts
warnings	Structured warnings, if any
metadata	Thread / message / timestamp

Annotated sample of a passing "answered" output (the text after `//` is an explanation, not part of the JSON):

```
{
  "schemaVersion": "trueblue.chat.output.v1",    // must be exactly this
  "templateId": "rag_qa.default.v1",             // the output template us
  "templateVersion": 1,
  "status": "answered",                          // a supported question
  "responseText": "The filing status is Single [S1].", // matches the or
  "sources": [                                   // only sources actually
    {
      "sourceId": "S1",
      "marker": "[S1]",
      "documentId": "doc_1",                     // the selected document
      "pageStart": 1,
      "pageEnd": 1,
      "pageLabel": "Page 1",
      "snippet": "Filing status: Single"          // the supporting text
    }
  ],
  "coverage": { "selectedDocumentIds": ["doc_1"], "noEvidenceDocumentIds": "" },
  "support": {
    "confidenceLabel": "medium",
    "retrievalMode": "vector_retrieval"          // PASS – the document i
  },
  "warnings": [],
}
```

```
"metadata": { "threadId": "thread_1", "messageId": "message_1", "respor
}
```

A PASS for Criterion 2 is: a valid `trueblue.chat.output.v1` envelope, with `responseText` matching the visible answer and `sources[]` listing the cited source(s).

 Structured output – Raw view

`support.confidenceLabel` is **source-support** confidence (how well retrieved evidence backs the answer), **not** tax-advice confidence. See Limitation 7.

TC-03 / TC-04 – Insufficient-evidence and non-document behavior

- **TC-03:** with a source selected, ask an out-of-document question (e.g. a spacecraft purchase). Expect `status = insufficient_evidence` – the system's honest "I could not find enough support in your documents to answer that" result – with **no citations** and a warning `INSUFFICIENT_EVIDENCE`.

 Insufficient evidence response

- **TC-04:** ask `hi`. Expect `status = non_document`, brief guidance, and no citations.

TC-05 – Configurable output templates ("client-defined schemas") – *developer-run*

This check is run by the **developer** and the result reported to you; **you are not expected to run any API commands**. The developer requests the compact registered template and then an unregistered template id, and reports:

- The compact request returns `output.templateId = rag_qa.compact.v1` with valid cited sources.
- An unregistered/unknown template id is rejected with **HTTP 400**.

Together these confirm that "client-defined schemas" are delivered as **approved, registered templates** rather than arbitrary runtime JSON (Limitation 2). You acknowledge this developer-run result in Sign-off.

TC-06 / TC-07 – Chat UI and thread replay

- Ask a supported question; confirm the streamed answer, inline `[S1]`, an expandable citation (filename + page + snippet), the **Insights** view (status/support/coverage), and the **Raw / Structured JSON** view.

- Confirm a new entry appears in **History (Chats)**; reopen it and confirm **the saved answer reopens with its citations and structured details intact**, with the source scope locked.

TC-08 – Source selection and multi-source honesty

Start a new thread, select 2–3 completed PDFs (include the **Whittaker** return), and ask:

For each selected return, what taxpayer name is shown?

Each selected return should be addressed — **the Whittaker return should return Jordan Whittaker**. Any return without supporting evidence for the asked field should be explicitly flagged (not silently omitted); citations should map to the documents actually used.

TC-09 – Tenant isolation spot-check

Confirm Acme data as `admin@acmetax.com`, then sign in as `admin@besttax.com` and confirm Acme documents and chats are **not** visible.

Note: data-layer tenant isolation holds — one firm cannot see or open another firm's documents. A separate **prompt-premise** wording edge case is tracked separately and described in **Limitation 10**; it is a wording/hygiene issue, **not** a cross-tenant data leak.

TC-10 – Deployed and accessible

Open `http://52.70.0.80/login`, log in, and open `/dashboard/chat`. The login page should load, login should succeed, and the dashboard and chat should load without server errors.

TC-11 – Deployment documentation — *developer-attested*

No client action is required. The developer attests that the deployment documentation delivered with this handoff (deployment notes, the output-contract reference, and this guide) is sufficient for an engineer to redeploy the system and re-run M4 validation independently — including environment variables, the database migration, and validation steps. You acknowledge this attestation in Sign-off; you do **not** need to open any repository files.

Scope and Limitations

These are stated plainly so acceptance is honest. They define the boundary between a **controlled M4 demo on non-sensitive samples** (supported now) and **unrestricted real-tax-document use** (not yet).

1. **HTTP only – use redacted/non-sensitive samples only.** Staging is served over plain HTTP. Passwords and uploads are unencrypted in transit. Do **not** upload real client tax data until TLS/HTTPS is added or the risk is explicitly waived.
2. **"Client-defined schemas" = approved registered templates, not arbitrary runtime schemas.** M4 delivers a configurable template engine with registered templates (`rag_qa.default.v1` , `rag_qa.compact.v1`). Unknown/unregistered templates are rejected (HTTP 400). New client output shapes are added by registering a template, not by posting an arbitrary JSON schema at runtime.
3. **Tax-domain advisory output is M6, not M4.** M4 proves the **source-cited structured envelope** (`responseText` + `citations` + `coverage` + retrieval-support `confidence` + `warnings`). The full 8-section tax advisory report, strategy categorization, and advice disclaimers are Milestone 6.
4. **Upload size: UI/proxy allows up to 25 MB, but the API rejects files over 20 MB.** A file between 20–25 MB can pass the proxy and then be rejected by the application (HTTP 413, "exceeds maximum of 20MB"). Recommendation: align the UI/proxy limit to 20 MB. For testing, keep sample PDFs comfortably under 20 MB (all provided samples are well under).
5. **Uploads are serialized per app process.** Simultaneous uploads from multiple testers can return `429 Too Many Requests` . Upload one document at a time and wait for `COMPLETED` .
6. **Operational docs/tooling still reference some legacy paths.** These are being cleaned up before a polished, single-source redeploy runbook is finalized. They do not affect client testing but are noted for transparency.
7. **PII redaction is not yet implemented, and numeric field-mapping is pending SME certification.**
 - The chat can echo identifier PII (SSN, bank/routing/account numbers, email, phone) verbatim from a source document. There is no output-layer masking yet. This is the primary reason to use **redacted/non-sensitive samples only** until a redaction pass ships.
 - The system enforces that answers are **grounded** and tied to a **real page** of a real document, and it strips unsupported numbers. It does **not** yet certify that a given figure maps to the **intended tax line** (e.g. that "9,649" is the *total tax* line). Numeric **field-mapping correctness** requires subject-matter-expert (SME) certification, which is pending.
8. **Response time – expect a few seconds per answer.** Each response is retrieved, re-ranked for relevance, and grounded before it is returned, so answers typically take **a few seconds** (up to ~10–15s for complex multi-document questions, or for the first request after an idle period). This is expected behaviour, not an error.
9. **Relevance re-ranking runs on a shared service tier.** Retrieved evidence is re-ranked for relevance before answering. Under heavy **simultaneous** testing this may slow responses or briefly fall back to standard retrieval ordering – the system **degrades gracefully and does not fail**. For the cleanest, most comparable results, have testers run **sequentially** where practical.
10. **Deferred prompt-hardening.** A deliberately misleading prompt – for example, "list another firm's clients" – can cause the assistant to *mislabel the user's own firm data*. This is a wording/hygiene issue, **not** a cross-tenant data leak: data-layer tenant isolation still holds (foreign-document probes fail closed). A prompt-hardening fix is **deferred**.

Additional product notes (consistent with prior milestones):

- AI wording may vary between runs, but answers must remain supported by the selected source evidence.
 - `insufficient_evidence` is a **correct** result when the selected documents do not support the question.
 - OCR for scanned/image PDFs is M5; the full tax advisory categorization is M6. Neither should block M4 acceptance.
-

Acceptance Criteria Checklist

Use this for sign-off. Each item is either **client-checkable** in the browser or marked **(developer)** where it is confirmed by the developer and acknowledged in Sign-off.

Criterion 1 – End-to-end upload → process → query → structured response

- ☐ A text-based PDF uploads and reaches `COMPLETED` (or is already present as a selectable source).
- ☐ The document appears as a selectable source.
- ☐ A supported query returns a grounded answer with a citation (for the Whittaker return: filing status = `Single` ; wages / total income = `$27,645`).
- ☐ The Raw / Structured JSON view shows `output.schemaVersion = "trueblue.chat.output.v1"` , `status = "answered"` , and `support.retrievalMode = "vector_retrieval"` .
- ☐ The cited source filename matches the document you selected. (Optional developer/API note: `output.sources[].documentId` equals the uploaded document's id.)

Criterion 2 – Output conforms to the agreed JSON schema

- ☐ `schemaVersion` , `templateId` , `templateVersion` , `status` , `responseText` , `sources` , `coverage` , `support` , `warnings` , `metadata` all present in the Raw view.
- ☐ `responseText` matches the visible answer; `sources[]` are final cited sources only.
- ☐ Unsupported question → `insufficient_evidence` , 0 sources, `INSUFFICIENT_EVIDENCE` warning.
- ☐ Greeting / non-document message → `non_document` , 0 sources.
- ☐ Compact template → `templateId = "rag_qa.compact.v1"` ; unknown template → HTTP 400. (**confirmed by developer – TC-05**)

Criterion 3 – Chat interface

- ☐ User can submit a query and see the answer in the conversation.

- ☐ Inline citation marker + expandable citation (filename, page, snippet) appear for grounded answers.
- ☐ Insights view shows status/support/coverage; Raw / Structured JSON view shows the full `output`.
- ☐ History saves the thread; reopening replays the answer with its citations and structured details intact; source scope stays locked.
- ☐ Multi-source query addresses each selected document or flags no-evidence per document.

Criterion 4 – Deployed and accessible

- ☐ `http://52.70.0.80/login` loads.
- ☐ Login succeeds; `/dashboard/chat` loads without server errors.

Criterion 5 – Deployment documentation (developer-attested)

- ☐ The developer attests that the deployment documentation delivered with this handoff is sufficient for an engineer to redeploy and re-run M4 validation independently. **Client acknowledges in Sign-off** (no repository files need to be opened).

Supporting checks

- ☐ Best Tax cannot see/access Acme documents or chats. (See *Limitation 10* for the deferred prompt-premise wording item.)
- ☐ In the Raw view, the public output shows **no system-prompt text, keys, or secrets**. (If anything like that appears, that is a FAIL – report it to the developer.)

Sign-off

Field	Value
Milestone	M4 – Output Generation, Testing & Deployment
Build tag	<code>m4-20260627-191552-ed97574</code> (app + worker)
Environment	<code>http://52.70.0.80</code> (HTTP, staging), AWS account <code>536573256060</code>
Schema version	<code>trueblue.chat.output.v1</code>
Developer attestation	The developer attests M4 was validated end-to-end on staging; the configurable-template / registered-schema check (TC-05) and the deployment documentation (TC-11) are confirmed by the developer.
Acceptance scope	Controlled M4 testing on non-sensitive / redacted sample PDFs over HTTP

Sign-off acknowledges: the five SOW acceptance criteria above are met on staging for non-sensitive sample documents – including the **developer-attested items** (the configurable-template behavior in TC-05 and the deployment documentation in TC-11) – **and** that the Scope and Limitations section (especially HTTP-only, no PII redaction, and pending SME numeric-field certification) is understood. Sign-off does **not** authorize unrestricted real-client-tax-document use; that is gated on TLS/HTTPS + an output-layer PII redaction pass (or an explicit written waiver).

Role	Name	Decision (Accept / Accept-with-conditions / Reject)	Date
Client approver (EulerTel)			
Delivery (developer)			

Conditions / notes:

(record any conditions, deferred items accepted, or follow-ups here)